

# Lab: K-Means

Rose Hoover



## Getting Started

1. Download the `.qmd` file from Moodle and any needed `.xlsx` or `.csv` data files. Save these in the *same folder/directory*.
2. Open the Quarto file in RStudio: **File > Open File... >**. If you're working on the MHC RStudio server, you need to upload the files first: go to the **Files** panel, then click **Upload**. Upload the `.qmd` file and any data files. You will need to upload each file one at a time.
3. Update the author and date in the YAML header of this file.
4. Click the **Render** button. If successful, you should have a new window pop up with a nice looking HTML document.
5. For this lab, you **may need** to still the package `glmnet`.

**Ask for help** if you encounter issues on any of the steps above. Once you've successfully made it through these steps, you can continue.

## Load Packages

You likely will need to install some these packages before you can run the code chunk below successfully.

```
library(tidyverse)
library(palmerpenguins)
library(factoextra)
library(ama)
```

## Load Penguin Data

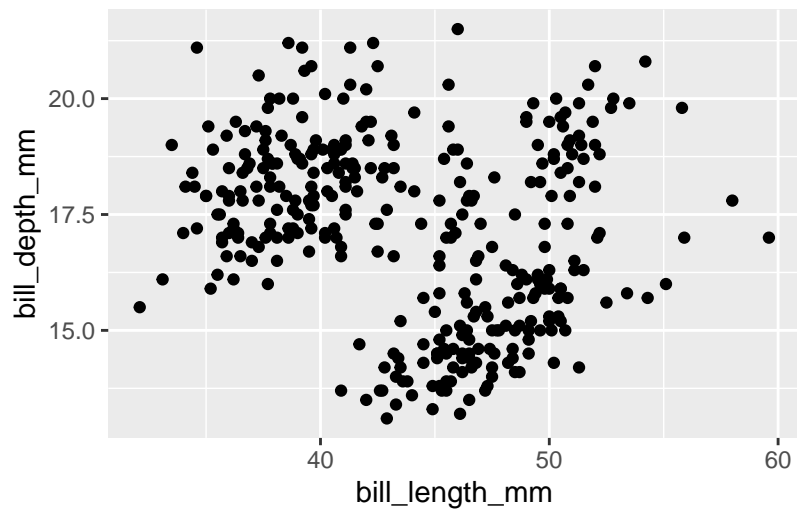
```
data(penguins)
```

## Data Cleaning

```
# Remove missing values
# YOUR CODE HERE
penguins = penguins %>%
  filter(!is.na(bill_length_mm) & !is.na(bill_depth_mm) & !is.na(species))
# Make data table (named penguins_reduced) that only has
# bill_length_mm and bill_depth_mm columns
penguins_reduced <- penguins %>% select(bill_length_mm, bill_depth_mm)
```

## Initial Visualization

```
ggplot(penguins, aes(x = bill_length_mm, y = bill_depth_mm)) +
  geom_point()
```



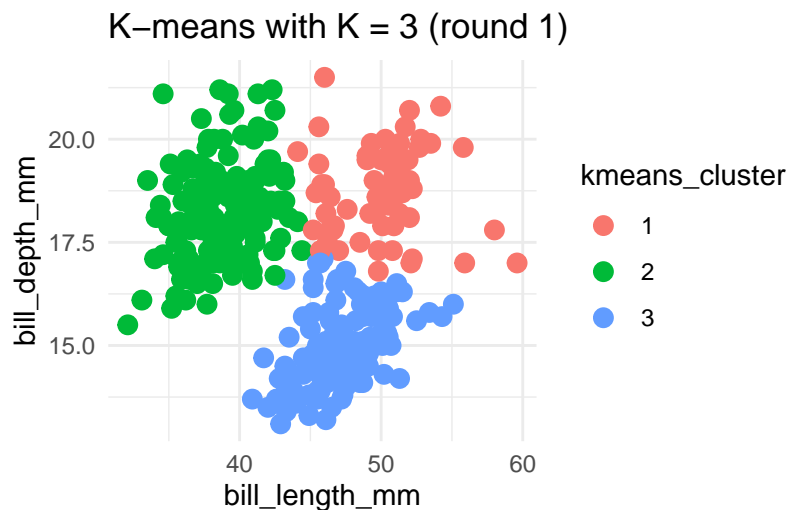
We'll cluster these penguins based on their bill lengths and depths:

## Implement $K$ -Means

Complete the code below to run the K-means algorithm using  $K = 3$ .

```
set.seed(244)
# Run the K-means algorithm
kmeans_3_round_1 <- kmeans(scale(penguins_reduced), centers = 3)

# Plot the cluster assignments
penguins_reduced %>%
  mutate(kmeans_cluster = as.factor(kmeans_3_round_1$cluster)) %>%
  ggplot(aes(x = bill_length_mm, y = bill_depth_mm, color = kmeans_cluster))
  ↪ +
  geom_point(size = 3) +
  theme(legend.position = "none") +
  labs(title = "K-means with K = 3 (round 1)") +
  theme_minimal()
```

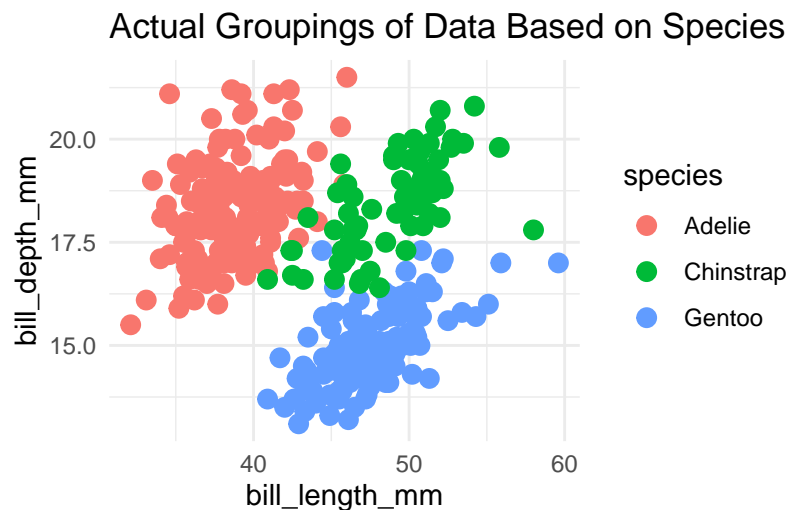


- Why do we have to set the seed for K-means? In practice, why should we try out a variety of seeds?

**Answer.** K means is a greedy algorithm and its possible so its possible for some random locations to have better/different results each time. It is also possible for it to get stuck at a local solution

## K-Means Clusters Versus Known Species Groupings

```
ggplot(data = penguins, aes(x = bill_length_mm, y = bill_depth_mm, color =  
  ↪ species)) +  
geom_point(size = 3) +  
theme(legend.position = "none") +  
labs(title = "Actual Groupings of Data Based on Species") +  
theme_minimal()
```



- Visually, how well do you think  $K$ -means captured the underlying species structure of the data?

**Answer.** It seems like it found roughly the data points corresponding to each of the species groups. It is identifying that as a way to structure the data.

## Tuning $K$

- To implement  $K$ -means clustering we must choose an appropriate  $K$ ! Use the following example to see the two different extreme situations. Typically, the ideal  $K$  is somewhere between the two extremes.
- **Minimum:**  $K = 2$  groups/clusters
- **Maximum:**  $K = n$  groups/clusters (one observation per cluster)

What happens in the  $K$ -means algorithm if  $K = n$ ?

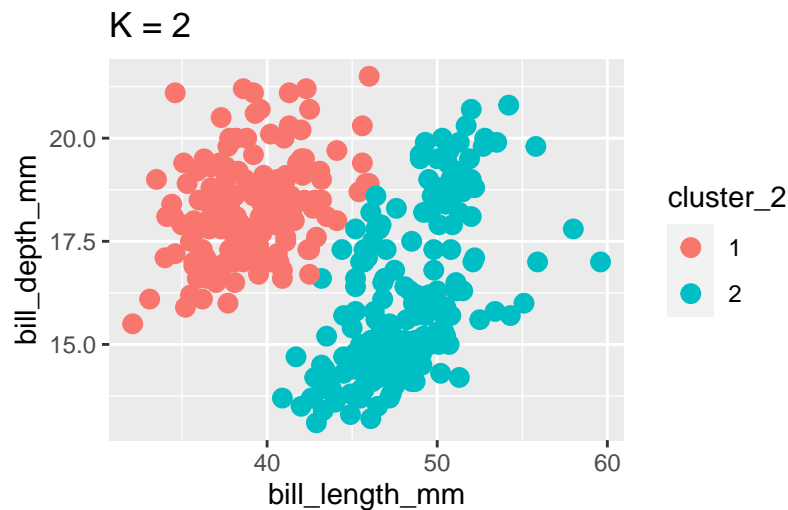
**Answer.** YOUR ANSWER HERE

Let's consider anywhere from  $K = 2$  to  $K = 20$  clusters.

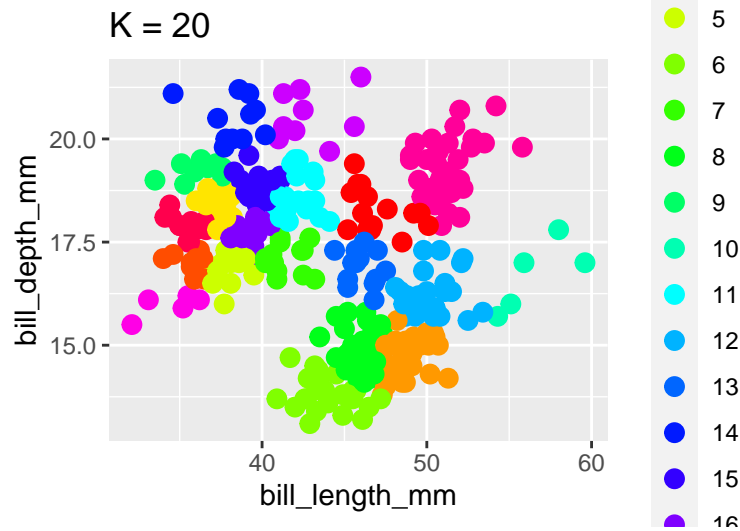
```
set.seed(244)

k_2 <- kmeans(scale(penguins_reduced), centers = 2)
k_20 <- kmeans(scale(penguins_reduced), centers = 20)

penguins_reduced %>%
  mutate(cluster_2 = as.factor(k_2$cluster)) %>%
  ggplot(aes(x = bill_length_mm, y = bill_depth_mm, color = cluster_2)) +
    geom_point(size = 3) +
    labs(title = "K = 2")
```



```
penguins_reduced %>%
  mutate(cluster_20 = as.factor(k_20$cluster)) %>%
  ggplot(aes(x = bill_length_mm, y = bill_depth_mm, color = cluster_20)) +
    geom_point(size = 3) +
    labs(title = "K = 20") +
    scale_color_manual(values = rainbow(20))
```



What are your general impressions?

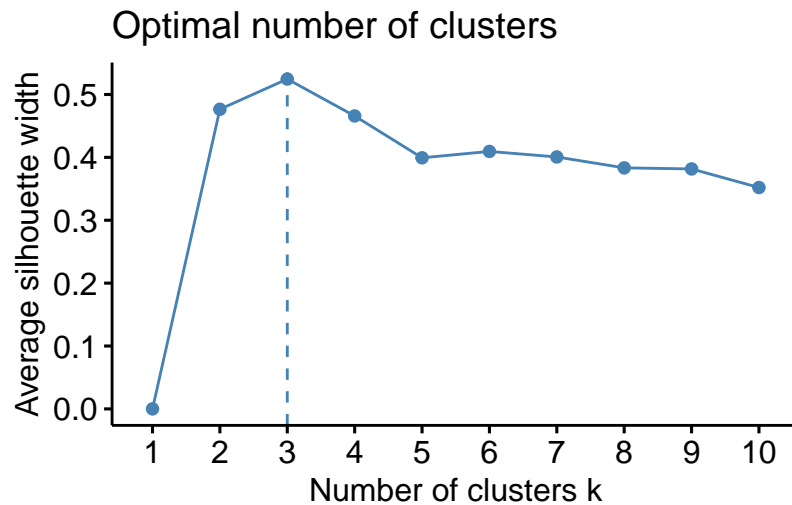
**Answer.** YOUR ANSWER HERE

## Finding Ideal K Value: Silhouette

- The *average silhouette approach* measures the quality of a clustering. That is, it determines how well each object lies within its cluster.
  - To do so, it maximizes the distance between clusters and minimizes distance within clusters.
- A high average silhouette indicates a good clustering.
- Given a range of possible K values, the optimal number of clusters (K) is the one that maximizes the average silhouette.

We can use a built-in silhouette method in the `fviz_nbclust` function to compute the average silhouette for various K values.

```
fviz_nbclust(scale(penguins_reduced), kmeans, method='silhouette')
```



Based on the *average silhouette approach*, what is the optimal  $K$  value?

**Answer.** optimal value is about 3!

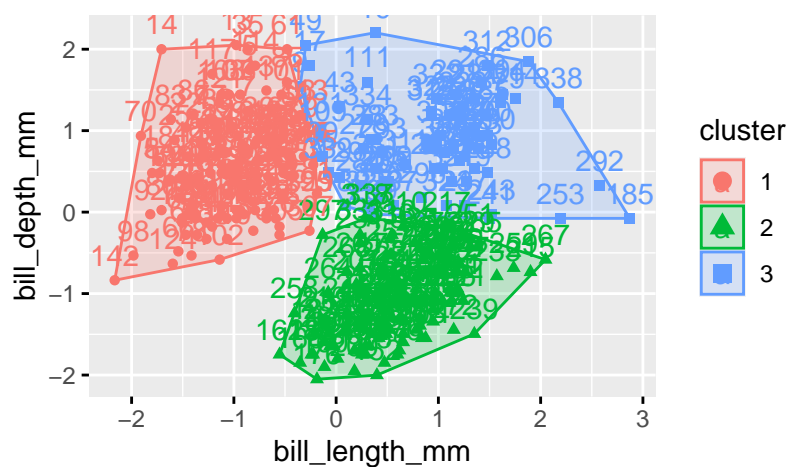
## Experimenting with Distance Metrics

We can use the `Kmeans` method (notice the “K” is capitalized in this function name) from the `amap` library to specify how we are measuring distance in the K-means algorithm.

```
set.seed(244)
k_2_manattan = Kmeans(scale(penguins_reduced), centers = 3,
                      method = "manhattan")
k_2_euclid = Kmeans(scale(penguins_reduced), centers = 3,
                    method = "euclidean")
k_2_maxnorm = Kmeans(scale(penguins_reduced), centers = 3,
                     method = "maximum")

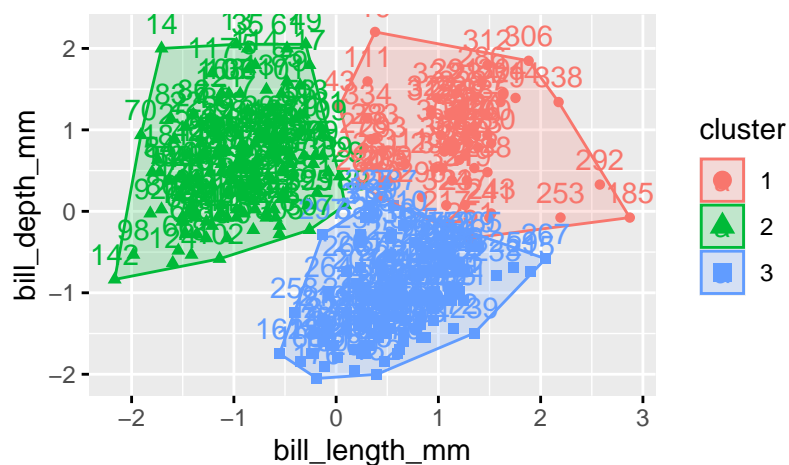
fviz_cluster(k_2_euclid, data = scale(penguins_reduced),
             main = sprintf("K = %d Clusters w/ Manhattan Distance", 3))
```

K = 3 Clusters w/ Manhattan Distance



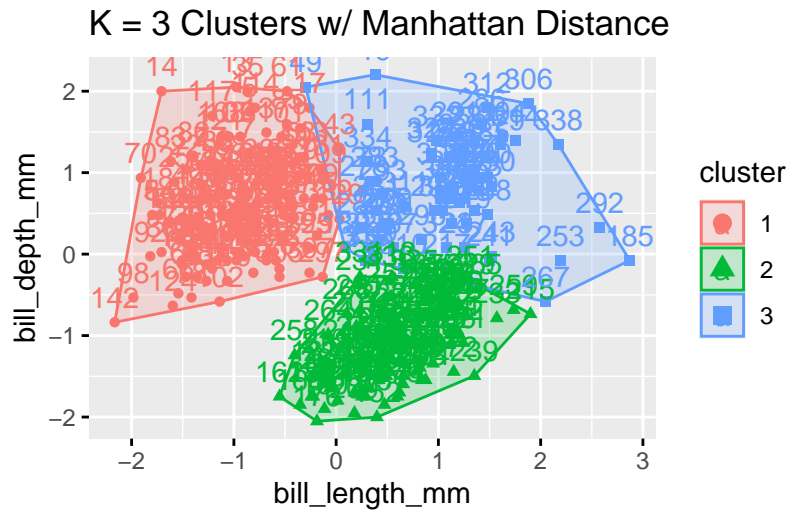
```
fviz_cluster(k_2_manattan, data = scale(penguins_reduced),
             main = sprintf("K = %d Clusters w/ Manhattan Distance", 3))
```

K = 3 Clusters w/ Manhattan Distance



```
fviz_cluster(k_2_maxnorm, data = scale(penguins_reduced),
             main = sprintf("K = %d Clusters w/ Manhattan Distance", 3))
```





Try changing  $K$  to equal 3 in the code chunk above. How do the clusterings using the 3 distance metrics compare? What do you generally observe?

**Answer.** YOUR ANSWER HERE

Modify the code in the chunk above so that we can easily change the value of  $K$  (rather than making sure to change  $K$  manually in every line). In general coding practices, is called *extracting out a constant*.